
NVIDIA GPUs: Detailed Reference Document

1. Company Overview

NVIDIA Corporation is an American technology company founded in 1993. The company designs GPUs and system-on-chip units. NVIDIA does not manufacture chips directly; it designs chips that are manufactured by semiconductor foundries.

NVIDIA GPUs are used in consumer, professional, and enterprise environments.

2. Definition of GPU

A GPU, or Graphics Processing Unit, is a specialized electronic circuit designed to rapidly manipulate and alter memory to accelerate the creation of images and data processing.

GPUs are designed for **parallel processing**. Parallel processing means performing many calculations at the same time.

3. CPU vs GPU

A CPU typically has:

- A small number of powerful cores
- Optimized for sequential tasks

A GPU typically has:

- Thousands of smaller cores
- Optimized for parallel tasks

GPUs are better suited for workloads such as deep learning and image processing.

4. NVIDIA GPU Core Components

4.1 CUDA Cores

CUDA cores are processing units inside NVIDIA GPUs. They execute arithmetic and logical operations.

The number of CUDA cores varies depending on the GPU model.

4.2 Streaming Multiprocessors

Streaming Multiprocessors, also known as SMs, are groups of CUDA cores.

Each SM contains:

- CUDA cores
 - Registers
 - Shared memory
 - Scheduling units
-

4.3 Registers

Registers are the fastest memory available on the GPU. Registers store temporary variables during computation.

4.4 Shared Memory

Shared memory is a small, fast memory shared between CUDA cores within the same SM.

Shared memory reduces access time compared to global memory.

4.5 Global Memory

Global memory is the largest memory space available to the GPU.

Global memory has higher latency compared to shared memory and registers.

5. Memory Types in NVIDIA GPUs

NVIDIA GPUs use multiple memory types:

- Registers
- Shared memory
- L1 cache
- L2 cache
- Global memory

Each memory type has different speed and size characteristics.

6. CUDA Programming Model

CUDA is NVIDIA's parallel computing platform.

CUDA programs use:

- Kernels
- Threads

- Blocks
- Grids

6.1 Threads

A thread is the smallest unit of execution in CUDA.

6.2 Blocks

A block is a group of threads.

Threads in the same block can share memory and synchronize.

6.3 Grids

A grid is a collection of blocks.

Each kernel launch creates a grid.

7. Tensor Cores – Detailed Explanation

Tensor Cores are specialized units in NVIDIA GPUs.

Tensor Cores accelerate:

- Matrix multiplication
- Mixed-precision arithmetic

Tensor Cores are commonly used in deep learning.

8. Precision Types Supported

NVIDIA GPUs support multiple precision types:

- FP32 (32-bit floating point)
- FP16 (16-bit floating point)
- BF16 (Brain Floating Point)
- INT8
- INT4

Lower precision improves performance at the cost of numerical accuracy.

9. NVIDIA GPU Architectures

NVIDIA GPU architectures are released in generations.

Common architectures include:

- Pascal
- Volta
- Turing
- Ampere
- Hopper

Each architecture improves performance and efficiency.

10. NVIDIA GPU Product Families

10.1 GeForce Series

GeForce GPUs are designed for gaming and consumer use.

10.2 RTX Series

RTX GPUs support real-time ray tracing.

RTX GPUs include Tensor Cores.

10.3 Data Center GPUs

Data center GPUs are designed for enterprise workloads.

Examples include:

- NVIDIA A100
 - NVIDIA H100
-

11. NVIDIA GPUs in Machine Learning

NVIDIA GPUs accelerate machine learning using CUDA and Tensor Cores.

Machine learning workloads include:

- Model training
- Model inference

Frameworks such as PyTorch and TensorFlow support NVIDIA GPUs.

12. NVIDIA GPUs in Data Centers

Data centers use NVIDIA GPUs for:

- AI workloads
- High-performance computing
- Data analytics

NVIDIA data center GPUs are optimized for reliability and performance.

13. Power and Thermal Considerations

NVIDIA GPUs consume significant electrical power.

Higher-end GPUs require:

- Adequate power supply
- Cooling systems

Thermal management is essential for stable performance.

14. Software Ecosystem

NVIDIA provides a software ecosystem that includes:

- CUDA Toolkit
- cuDNN
- TensorRT
- NVIDIA drivers

These tools improve performance and usability.

15. Virtualization and Multi-Instance GPU (MIG)

Some NVIDIA GPUs support Multi-Instance GPU technology.

MIG allows a single GPU to be partitioned into multiple isolated instances.

16. Limitations of NVIDIA GPUs

- High acquisition cost
 - Power consumption
 - Vendor lock-in with CUDA
-

17. Security and Isolation

NVIDIA GPUs support isolation features for enterprise use.

Isolation improves security in shared environments.

18. Summary

NVIDIA GPUs are parallel processors designed for high-performance workloads. They are widely used in gaming, artificial intelligence, and data centers. NVIDIA's architecture, CUDA platform, and Tensor Cores make these GPUs suitable for modern computing tasks.

19. Glossary

- **GPU:** Graphics Processing Unit
 - **CUDA:** Compute Unified Device Architecture
 - **SM:** Streaming Multiprocessor
 - **Tensor Core:** Specialized matrix computation unit
 - **Global Memory:** Main GPU memory
-